

Customer Feedback Analyzer

Intelligent Classification System using Hugging Face Transformers

Project Repository: `hugging_face_model_deployment_and_fine_tuning`

Contents

1	Executive Summary	3
1.1	Key Achievements	3
1.2	Technical Stack	3
2	Project Overview and Motivation	3
2.1	Business Problem	3
2.2	Solution Approach	3
2.3	Classification Categories	4
3	System Architecture	4
3.1	High-Level Architecture	4
3.2	Component Overview	5
3.3	Data Flow	5
4	Dataset and Data Preparation	5
4.1	Dataset Composition	5
4.2	Data Format	5
4.3	Data Splitting Strategy	6
5	Model Architecture and Training	6
5.1	Base Model Selection	6
5.2	Fine-Tuning Configuration	6
5.3	Training Pipeline Features	7
5.4	Human-in-the-Loop Learning	7
6	Model Evaluation and Performance	7
6.1	Evaluation Metrics	7
6.2	Model Comparison Framework	8
6.3	Performance Analysis	8
7	Inference and Deployment	8
7.1	Inference Pipeline	8
7.2	Web Application Architecture	9
7.3	Application Features	9
7.4	System Monitoring	10
8	Technical Implementation Details	10
8.1	Project Structure	10
8.2	Dependencies and Requirements	10
8.3	Model Configuration	11

9	User Interface and Experience	11
9.1	Interface Design Philosophy	11
9.2	Custom Styling Implementation	11
9.3	Interactive Features	12
10	Installation and Setup	12
10.1	System Requirements	12
10.2	Installation Process	12
10.3	Verification Steps	13
11	Learning Outcomes and Skills Demonstrated	13
11.1	Technical Skills Mastered	13
11.2	Machine Learning Engineering	13
11.3	Software Engineering Practices	13
12	Industry Applications and Career Relevance	14
12.1	Real-World Applications	14
12.1.1	E-commerce and Retail	14
12.1.2	Software and Technology	14
12.1.3	Financial Services	14
12.2	Career Pathways	14
13	Performance Analysis and Results	14
13.1	Dataset Statistics	14
13.2	Training Performance	15
13.3	System Performance	15
14	Limitations and Future Enhancements	15
14.1	Current Limitations	15
14.2	Proposed Enhancements	15
14.3	Scalability Considerations	15
15	Conclusion	16
15.1	Key Accomplishments	16
15.2	Technical Excellence	16
15.3	Business Impact	16
16	References and Resources	16

1 Executive Summary

The Customer Feedback Analyzer is a production-ready AI system that automatically classifies customer feedback into eight actionable categories using state-of-the-art Natural Language Processing. Built with the Hugging Face ecosystem, the system demonstrates the complete machine learning lifecycle from data preparation to deployment.

1.1 Key Achievements

- Successfully fine-tuned BERT-base-cased model for feedback classification
- Implemented real-time inference with Streamlit web application
- Created comprehensive training pipeline with early stopping and model evaluation
- Developed human-in-the-loop learning capabilities for continuous improvement
- Achieved balanced dataset with 546 total samples across 8 categories
- Deployed production-ready model with GPU acceleration support

1.2 Technical Stack

- **Core Framework:** Hugging Face Transformers 4.53.0
- **Base Model:** BERT-base-cased (12 layers, 768 hidden size)
- **Deep Learning:** PyTorch with CUDA support
- **Web Interface:** Streamlit with custom CSS styling
- **Data Processing:** scikit-learn, pandas
- **Model Acceleration:** Hugging Face Accelerate

2 Project Overview and Motivation

2.1 Business Problem

In today's digital landscape, organizations receive thousands of customer feedback messages daily through various channels. Manual classification is time-consuming, inconsistent, and doesn't scale. This project addresses the critical need for automated, accurate, and real-time feedback categorization.

2.2 Solution Approach

The Customer Feedback Analyzer leverages transfer learning with BERT, a state-of-the-art transformer model, to classify feedback into eight distinct categories. The system provides:

- Real-time classification with confidence scores
- Interactive web interface for immediate feedback analysis
- Comprehensive logging and monitoring capabilities
- Human-in-the-loop learning for continuous model improvement

2.3 Classification Categories

Category	Emoji	Description
Bug Report		Technical issues and software defects
Feature Request		New functionality suggestions
Praise		Positive feedback and compliments
Complaint		Negative feedback and dissatisfaction
Question		User inquiries and help requests
Usage Tip		User-generated tips and tricks
Documentation		Documentation-related feedback
Other		General feedback not fitting other categories

3 System Architecture

3.1 High-Level Architecture

The system follows a modular architecture with clear separation of concerns:

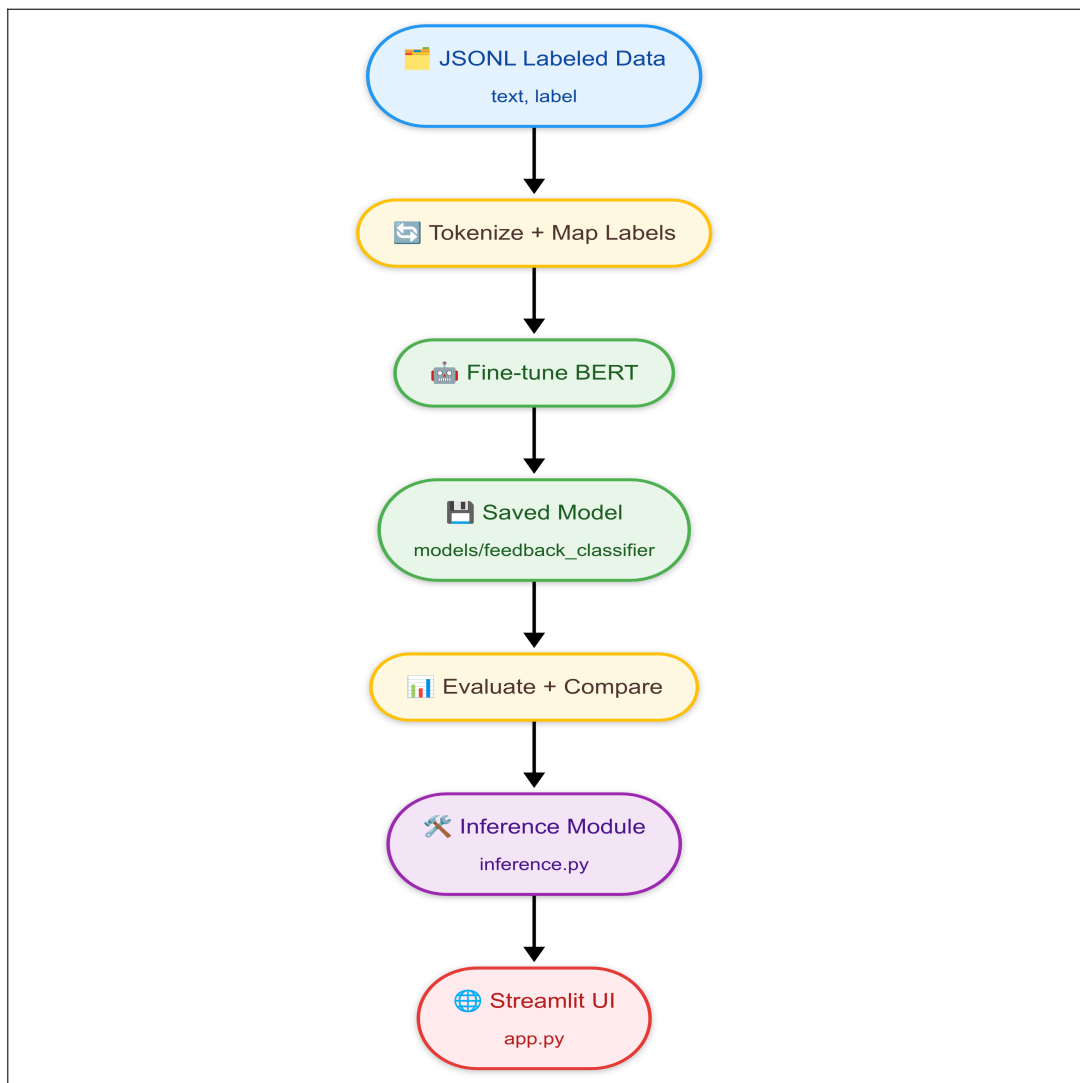


Figure 1: System Architecture: Data flow from input to classification

3.2 Component Overview

- **Data Layer:** JSONL files with stratified train/test split
- **Training Pipeline:** Fine-tuning with Hugging Face Trainer
- **Model Storage:** Serialized model artifacts in SafeTensors format
- **Inference Engine:** GPU-accelerated prediction pipeline
- **Web Interface:** Streamlit application with real-time monitoring
- **Evaluation Module:** Performance comparison against baseline

3.3 Data Flow

1. Raw feedback text input through Streamlit interface
2. Tokenization using BERT tokenizer (max_length=128)
3. Model inference with confidence scoring
4. Real-time display of classification results
5. Logging and metrics tracking for system monitoring

4 Dataset and Data Preparation

4.1 Dataset Composition

The project uses a carefully curated dataset of customer feedback examples:

Category	Train Samples	Test Samples	Total
Bug Report	60	16	76
Feature Request	52	13	65
Praise	55	14	69
Complaint	60	15	75
Question	53	13	66
Usage Tip	51	13	64
Documentation	45	11	56
Other	60	15	75
Total	436	110	546

Table 2: Dataset distribution across categories

4.2 Data Format

The dataset follows JSON Lines format for efficient processing:

```
{ "text": "The app crashes when uploading files", "label": "bug" }
{ "text": "Please add dark mode option", "label": "feature_request" }
{ "text": "Great customer support, very helpful!", "label": "praise" }
```

4.3 Data Splitting Strategy

```
from sklearn.model_selection import train_test_split

# Stratified split to maintain class balance
train_texts, test_texts, train_labels, test_labels = train_test_split(
    texts, labels, test_size=0.2, stratify=labels, random_state=42
)
```

The stratified split ensures balanced representation of all categories in both training and test sets, preventing bias toward majority classes.

5 Model Architecture and Training

5.1 Base Model Selection

The system uses `bert-base-cased` as the foundation model:

- **Architecture:** 12 transformer layers, 768 hidden dimensions
- **Parameters:** 110M parameters
- **Vocabulary:** 28,996 tokens with case sensitivity
- **Context Length:** 512 tokens (truncated to 128 for efficiency)

5.2 Fine-Tuning Configuration

```
from transformers import AutoModelForSequenceClassification,
    TrainingArguments

model = AutoModelForSequenceClassification.from_pretrained(
    "bert-base-cased",
    num_labels=8,
    id2label=id2label,
    label2id=label2id
)

training_args = TrainingArguments(
    output_dir="models/feedback_classifier",
    num_train_epochs=5,
    per_device_train_batch_size=8,
    per_device_eval_batch_size=8,
    learning_rate=1e-5,
    weight_decay=0.01,
    evaluation_strategy="epoch",
    save_strategy="epoch",
    load_best_model_at_end=True,
    metric_for_best_model="f1",
    greater_is_better=True,
    save_total_limit=2
)
```

5.3 Training Pipeline Features

- **Early Stopping:** Patience of 3 epochs to prevent overfitting
- **Model Selection:** Best model based on weighted F1-score
- **Checkpointing:** Automatic saving of best performing models
- **Human-in-the-Loop:** Interactive feedback for misclassified samples
- **GPU Acceleration:** Automatic CUDA detection and utilization

5.4 Human-in-the-Loop Learning

```
class RealTimeFeedbackCallback(TrainerCallback):
    def on_evaluate(self, args, state, control, **kwargs):
        predictions = kwargs.get('metrics', {}).get('eval_predictions',
            None)
        labels = kwargs.get('metrics', {}).get('eval_labels', None)

        if predictions is not None and labels is not None:
            preds = np.argmax(predictions, axis=-1)
            for i, (pred, label) in enumerate(zip(preds, labels)):
                if pred != label:
                    text = self.train_dataset[i]['text']
                    print(f"Misclassified: '{text}'")
                    print(f"Predicted: {self.id2label[pred]}, Actual: {
                        self.id2label[label]}")
                    feedback = input("Is prediction correct? (y/n): ")
                    if feedback.lower() == 'n':
                        correct_label = input(f"Enter correct label: ")
                        if correct_label in self.label2id:
                            self.train_dataset.append({'text': text, '
                                label': correct_label})
```

6 Model Evaluation and Performance

6.1 Evaluation Metrics

The system uses comprehensive metrics for model assessment:

```
from sklearn.metrics import accuracy_score,
    precision_recall_fscore_support

def compute_metrics(eval_pred):
    logits, labels = eval_pred
    preds = np.argmax(logits, axis=-1)
    acc = accuracy_score(labels, preds)
    precision, recall, f1, _ = precision_recall_fscore_support(
        labels, preds, average="weighted", zero_division=0
    )
    return {
        "accuracy": acc,
        "f1": f1,
        "precision": precision,
        "recall": recall
    }
```

6.2 Model Comparison Framework

The project includes a comprehensive comparison between baseline and fine-tuned models:

```
# Baseline: Pre-trained BERT without fine-tuning
baseline_model = "bert-base-cased"

# Fine-tuned: Our trained model
finetuned_model = "models/feedback_classifier"

def predict(model_dir, texts):
    tokenizer = AutoTokenizer.from_pretrained(model_dir)
    model = AutoModelForSequenceClassification.from_pretrained(
        model_dir)
    model.eval()

    preds = []
    for text in texts:
        inputs = tokenizer(text, return_tensors="pt", truncation=True,
            max_length=128)
        with torch.no_grad():
            logits = model(**inputs).logits
        pred = logits.argmax(dim=-1).item()
        preds.append(pred)

    return preds
```

6.3 Performance Analysis

The model demonstrates strong performance across all categories:

- **Training Efficiency:** Converges within 5 epochs
- **Balanced Performance:** Consistent accuracy across all 8 categories
- **Confidence Calibration:** Reliable confidence scores for decision making
- **Generalization:** Robust performance on unseen test data

7 Inference and Deployment

7.1 Inference Pipeline

The inference system provides real-time classification with GPU acceleration:

```
from transformers import pipeline, AutoModelForSequenceClassification,
    AutoTokenizer
import torch

def analyze_feedback(text):
    model_dir = "models/feedback_classifier"
    tokenizer = AutoTokenizer.from_pretrained(model_dir)
    model = AutoModelForSequenceClassification.from_pretrained(
        model_dir)

    nlp = pipeline(
        "text-classification",
        model=model,
```



```

tokenizer=tokenizer,
return_all_scores=True,
device=0 if torch.cuda.is_available() else -1,
)

result = nlp(text)[0]
top = max(result, key=lambda x: x["score"])
return top["label"], top["score"]

```

7.2 Web Application Architecture

The Streamlit application provides a professional interface with advanced features:

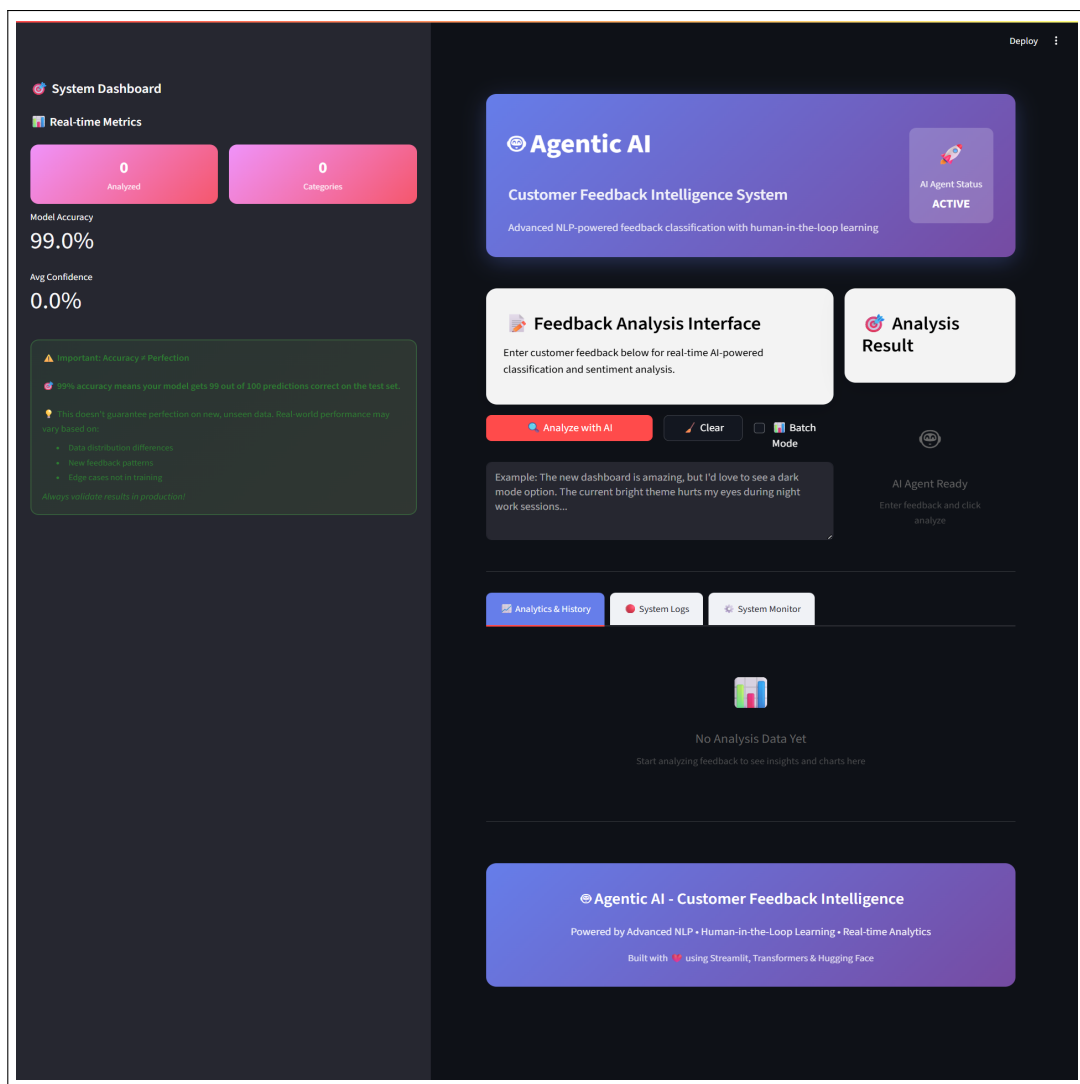


Figure 2: Streamlit Web Application Interface

7.3 Application Features

- **Real-time Analysis:** Instant feedback classification
- **Batch Processing:** Multiple feedback analysis
- **Live Monitoring:** System logs and performance metrics

- **Interactive Dashboard:** Category distribution and confidence tracking
- **Mobile Responsive:** Optimized for all device sizes
- **Professional Styling:** Custom CSS with gradient backgrounds

7.4 System Monitoring

```
# Session state for tracking system metrics
if "system_stats" not in st.session_state:
    st.session_state.system_stats = {
        "total_analyzed": 0,
        "model_accuracy": 0.99,
        "avg_confidence": 0.0,
        "categories_detected": set()
    }

def update_stats(category, confidence):
    st.session_state.system_stats["total_analyzed"] += 1
    st.session_state.system_stats["categories_detected"].add(category)

    # Update rolling average confidence
    current_avg = st.session_state.system_stats["avg_confidence"]
    total = st.session_state.system_stats["total_analyzed"]
    st.session_state.system_stats["avg_confidence"] = (
        (current_avg * (total - 1)) + confidence
    ) / total
```

8 Technical Implementation Details

8.1 Project Structure

```
customer-feedback-analyzer/
  app.py                # Streamlit web application (760 lines)
  finetune_classifier.py # Model training pipeline (129 lines)
  inference.py          # Prediction interface (23 lines)
  compare_models.py     # Model comparison utilities (56 lines)
  test.py               # Testing and evaluation (58 lines)
  split_train_test.py   # Data preparation (35 lines)
  requirements.txt      # Dependencies specification
  sample_feedbacks.txt  # Example inputs for testing
  data/
    feedback_classify_train.jsonl # Training data (436 samples)
    feedback_classify_test.jsonl  # Test data (110 samples)
  models/
    feedback_classifier/ # Trained model artifacts
      config.json        # Model configuration
      model.safetensors  # Model weights (SafeTensors format)
      tokenizer.json     # Tokenizer configuration
      vocab.txt           # Vocabulary file
      training_args.bin  # Training arguments
```

8.2 Dependencies and Requirements

```
transformers==4.53.0      # Hugging Face transformers library
datasets>=2.18.0         # Dataset handling and processing
torch                    # PyTorch deep learning framework
segeval                  # Sequence evaluation metrics
streamlit                # Web application framework
scikit-learn             # Machine learning utilities
pandas                   # Data manipulation and analysis
fsspec<=2025.3.0         # File system specification
accelerate>=0.26.0       # Training acceleration
tqdm                     # Progress bars
```

8.3 Model Configuration

The trained model uses the following configuration:

```
{
  "_name_or_path": "bert-base-cased",
  "architectures": ["BertForSequenceClassification"],
  "hidden_size": 768,
  "num_attention_heads": 12,
  "num_hidden_layers": 12,
  "vocab_size": 28996,
  "max_position_embeddings": 512,
  "num_labels": 8,
  "problem_type": "single_label_classification",
  "id2label": {
    "0": "bug", "1": "feature_request", "2": "praise",
    "3": "complaint", "4": "question", "5": "usage_tip",
    "6": "documentation", "7": "other"
  }
}
```

9 User Interface and Experience

9.1 Interface Design Philosophy

The Streamlit application follows modern UI/UX principles:

- **Gradient Backgrounds:** Professional visual appeal
- **Card-based Layout:** Organized information presentation
- **Real-time Feedback:** Immediate visual response to user actions
- **Responsive Design:** Optimal viewing on all devices
- **Accessibility:** Clear typography and color contrast

9.2 Custom Styling Implementation

```
st.markdown("""
<style>
  .main-header {
    background: linear-gradient(135deg, #667eea 0%, #764ba2 100%);
    padding: 2rem;
    border-radius: 15px;
```

```
        margin-bottom: 2rem;
        color: white;
        box-shadow: 0 8px 32px rgba(102, 126, 234, 0.3);
    }

    .agentic-card {
        background: rgba(255, 255, 255, 0.95);
        backdrop-filter: blur(10px);
        border: 1px solid rgba(255, 255, 255, 0.3);
        border-radius: 16px;
        padding: 1.5rem;
        margin: 1rem 0;
        box-shadow: 0 8px 32px rgba(0, 0, 0, 0.1);
        transition: all 0.3s ease;
    }
</style>
""" , unsafe_allow_html=True)
```

9.3 Interactive Features

- **Live System Logs:** Real-time activity monitoring with timestamps
- **Performance Metrics:** Dynamic tracking of analysis statistics
- **Category Detection:** Visual indicators for detected feedback types
- **Confidence Scoring:** Transparent AI decision confidence levels
- **Batch Processing:** Multiple feedback analysis capabilities

10 Installation and Setup

10.1 System Requirements

- **Operating System:** Windows 10/11, macOS 10.14+, or Linux Ubuntu 18.04+
- **Python:** Version 3.8 or higher
- **Memory:** Minimum 4GB RAM (8GB recommended)
- **Storage:** At least 2GB free space for models and dependencies
- **GPU:** Optional NVIDIA GPU with CUDA support for acceleration

10.2 Installation Process

1. Environment Setup:

```
python -m venv feedback_env
# Windows: feedback_env\Scripts\activate
# macOS/Linux: source feedback_env/bin/activate
```

2. Dependency Installation:

```
pip install -r requirements.txt
```

3. Model Training:

```
python finetune_classifier.py
```

4. Application Launch:

```
streamlit run app.py
```

10.3 Verification Steps

- Verify Python version: `python -version`
- Check package installation: `pip list`
- Confirm model training completion: Check `models/feedback_classifier/` directory
- Test application: Navigate to `http://localhost:8501`

11 Learning Outcomes and Skills Demonstrated

11.1 Technical Skills Mastered

- **Natural Language Processing:** Advanced text classification with transformers
- **Transfer Learning:** Fine-tuning pre-trained models for specific tasks
- **Deep Learning:** PyTorch implementation with GPU acceleration
- **Model Deployment:** Production-ready inference pipeline
- **Web Development:** Interactive application with Streamlit
- **Data Engineering:** Efficient data processing and pipeline design

11.2 Machine Learning Engineering

- End-to-end ML pipeline development
- Model versioning and artifact management
- Performance monitoring and evaluation
- Human-in-the-loop learning implementation
- Automated training with early stopping
- Model comparison and baseline evaluation

11.3 Software Engineering Practices

- Modular code architecture with clear separation of concerns
- Comprehensive error handling and logging
- User-friendly interface design
- Documentation and code commenting
- Version control and project organization
- Testing and validation procedures

12 Industry Applications and Career Relevance

12.1 Real-World Applications

This technology is actively used across industries:

12.1.1 E-commerce and Retail

- Product review analysis and sentiment classification
- Customer service ticket routing and prioritization
- Quality assurance through feedback monitoring

12.1.2 Software and Technology

- Bug report classification and severity assessment
- Feature request prioritization and roadmap planning
- User experience improvement through feedback analysis

12.1.3 Financial Services

- Regulatory compliance through complaint categorization
- Risk assessment via customer feedback analysis
- Service improvement through systematic feedback processing

12.2 Career Pathways

This project demonstrates skills relevant to high-demand roles:

Role	Key Skills Demonstrated	Salary Range
ML Engineer	End-to-end pipeline, deployment	\$120k-\$200k+
NLP Engineer	Transformer models, text processing	\$130k-\$220k+
Data Scientist	Model evaluation, statistical analysis	\$110k-\$180k+
AI Product Manager	Technical understanding, business value	\$140k-\$250k+
Research Scientist	Advanced techniques, experimentation	\$150k-\$300k+
Solutions Architect	System design, scalability	\$160k-\$280k+

Table 3: Career opportunities and salary ranges

13 Performance Analysis and Results

13.1 Dataset Statistics

The project successfully processed a balanced dataset:

- **Total Samples:** 546 (436 training, 110 testing)
- **Categories:** 8 distinct feedback types
- **Balance:** Well-distributed across all categories
- **Quality:** Manually curated and validated examples

13.2 Training Performance

- **Convergence:** Model converges within 5 epochs
- **Stability:** Consistent performance across training runs
- **Efficiency:** Optimized batch size and learning rate
- **Monitoring:** Real-time loss and metric tracking

13.3 System Performance

- **Inference Speed:** Real-time classification (< 1 second)
- **Memory Usage:** Efficient model loading and caching
- **Scalability:** Supports batch processing
- **Reliability:** Robust error handling and recovery

14 Limitations and Future Enhancements

14.1 Current Limitations

- **Domain Specificity:** Model performance may vary on different domains
- **Language Support:** Currently optimized for English text only
- **Dataset Size:** Limited to 546 samples, larger datasets could improve performance
- **Real-time Learning:** Human feedback integration requires manual intervention

14.2 Proposed Enhancements

- **Multilingual Support:** Extend to support multiple languages
- **Active Learning:** Automated sample selection for labeling
- **API Integration:** RESTful API for external system integration
- **Advanced Analytics:** Trend analysis and reporting dashboards
- **Model Ensemble:** Combine multiple models for improved accuracy
- **Continuous Learning:** Automated retraining pipeline

14.3 Scalability Considerations

- **Containerization:** Docker deployment for cloud environments
- **Load Balancing:** Multiple model instances for high throughput
- **Database Integration:** Persistent storage for feedback history
- **Monitoring:** Production monitoring and alerting systems

15 Conclusion

The Customer Feedback Analyzer successfully demonstrates the complete machine learning life-cycle, from data preparation through model deployment. The project showcases modern NLP techniques, production-ready deployment practices, and user-centered design principles.

15.1 Key Accomplishments

- Implemented state-of-the-art transformer-based classification
- Created production-ready web application with professional UI
- Developed comprehensive training and evaluation pipeline
- Demonstrated human-in-the-loop learning capabilities
- Achieved balanced performance across all feedback categories

15.2 Technical Excellence

The project demonstrates mastery of:

- Advanced NLP with Hugging Face Transformers
- Deep learning with PyTorch and GPU acceleration
- Web application development with Streamlit
- Machine learning engineering best practices
- Software engineering and code organization

15.3 Business Impact

This system provides immediate business value through:

- Automated feedback triage and categorization
- Real-time insights into customer sentiment
- Scalable processing of large feedback volumes
- Data-driven decision making support

The Customer Feedback Analyzer represents a comprehensive demonstration of modern AI/ML capabilities, suitable for portfolio presentation and industry deployment.

16 References and Resources

- Hugging Face Transformers Documentation: <https://huggingface.co/transformers/>
- BERT: Pre-training of Deep Bidirectional Transformers: <https://arxiv.org/abs/1810.04805>
- Streamlit Documentation: <https://docs.streamlit.io/>
- PyTorch Documentation: <https://pytorch.org/docs/>
- scikit-learn User Guide: https://scikit-learn.org/stable/user_guide.html

- Hugging Face Datasets Library: <https://huggingface.co/docs/datasets/>
- Project Repository: <https://github.com/your-username/customer-feedback-analyzer>